Sentinel Real-Time Deepfake Detection System: A Comprehensive Technical Analysis of Architecture, Algorithmic Methodologies, and Forensic Robustness

Executive Summary

The digital information ecosystem is currently besieged by an unprecedented proliferation of hyper-realistic synthetic media, colloquially known as "deepfakes." These manipulated assets, generated through increasingly sophisticated adversarial machine learning techniques, pose a systemic risk to democratic institutions, financial stability, and individual reputational integrity. The Sentinel Real-Time Deepfake Detection System emerges as a critical counter-measure in this adversarial landscape. This report provides an exhaustive technical dissection of the Sentinel architecture, a system designed to operate at the intersection of high-throughput real-time inference and forensic-grade accuracy.

Sentinel differentiates itself from first-generation detection systems by abandoning reliance on transient visual artifacts—such as blinking anomalies or warping—which generative models have successfully evolved to conceal. Instead, Sentinel adopts a multi-modal detection philosophy anchored in the **EfficientNet-B7** architecture, utilizing a **Noisy Student** semi-supervised training paradigm to achieve robustness against the "digital decay" of compression and transmission. Furthermore, the system integrates the **Face X-ray** methodology, a novel approach that targets the fundamental structural blending boundaries inherent to face-swapping algorithms, thereby granting the system a generalized detection capability against unseen forgery methods.

This analysis draws upon a wide array of academic literature, technical documentation, and empirical benchmarks to evaluate Sentinel's performance. It scrutinizes the system's reliance on the **Deepfake Detection Challenge (DFDC)** dataset for training, examining how the

diversity of this corpus influences model bias and generalization. We delve into the specific data augmentation strategies employed via the **Albumentations** library—techniques with a pedigree tracing back to winning solutions in the Kaggle 2018 Data Science Bowl—to understand how Sentinel simulates the hostile environment of the open web.

The report further explores the theoretical underpinnings of **frequency domain analysis**, elucidating how Sentinel implicitly leverages spectral discrepancies to identify GAN "fingerprints" that are invisible to the human eye. By synthesizing insights on neuron coverage monitoring (FakeSpotter) and the systematic flaws of Convolutional Neural Network (CNN) generators, we construct a holistic view of the system's detection logic. We conclude with a forward-looking assessment of the challenges posed by emerging diffusion-based generators and the necessity for multi-modal (audio-visual) coherence checking in future iterations.

1. The Deepfake Threat Landscape: Evolution and Taxonomy

1.1 The Genesis and Democratization of Synthetic Media

To understand the architectural decisions behind the Sentinel system, one must first appreciate the trajectory of the threat it is designed to neutralize. The term "deepfake" is a portmanteau of "deep learning" and "fake," originating from a Reddit user in 2017 who utilized open-source machine learning libraries to swap celebrity faces into pornographic content.[1] However, the underlying technologies—Autoencoders and Generative Adversarial Networks (GANs)—predate this cultural inflection point.

Early manipulation techniques were the domain of specialized visual effects studios, requiring manual rotoscoping and substantial compute farms. The revolution brought about by deep learning was the automation of this mapping process. The initial generation of deepfake tools, such as **DeepFake AutoEncoder (DFAE)**, utilized a shared encoder and two decoder networks. The encoder would compress the face of Person A into a latent representation, and the decoder for Person B would reconstruct it. When the face of Person A was fed into the encoder and then passed to Person B's decoder, the system would reconstruct Person A's expression with Person B's likeness.[2]

These early iterations were plagued by low resolution and obvious artifacts: lack of blinking, static gaze, and blurry boundaries where the face met the hairline. However, the introduction of **Generative Adversarial Networks (GANs)** accelerated the quality of synthesis. In a GAN framework, a generator network attempts to create a fake image, while a discriminator network attempts to distinguish it from a real image. The two networks train in a zero-sum game, driving the generator to produce increasingly realistic textures, lighting, and micro-expressions.[3]

Today, we face a third generation of synthesis: **Diffusion Models** and **Neural Radiance Fields (NeRFs)**. These technologies move beyond simple 2D pixel manipulation to model the underlying 3D geometry and lighting physics of the scene, producing "in-the-wild" deepfakes that exhibit temporal stability and high-frequency detail previously thought impossible to synthesize.[4]

1.2 Taxonomy of Facial Manipulation

The Sentinel system is not a monolith; its detection capabilities are tuned to specific categories of manipulation defined in the forensic literature. Understanding these categories is essential for evaluating the system's scope.

Manipulation Type	Description	Technical Mechanism	Sentinel's Primary Detection Vector
Identity Swap	Replacing the face of a source subject with a target subject.	Autoencoders, GANs (e.g., FaceSwap, DeepFaceLab).	**Face X-ray**: Detecting the blending boundary between the swapped face and the original background.[6]
Face Reenactment	Manipulating the expression or head pose of a source actor to drive a target.	Neural Talking Heads (NTH), First Order Motion Models.	**Temporal Inconsistency**: Detecting jitter or unnatural muscle movements over time.[7]
Attribute	Altering specific features (age,	StarGAN, AttGAN.	**Spectral Analysis**: Identifying

Manipulation	gender, hair, glasses) without changing identity.		high-frequency artifacts in the modified regions.[1]
Entire Face Synthesis	Generating non-existent identities from scratch.	StyleGAN, ProGAN.	**Global Consistency**: Checking for asymmetries (e.g., mismatched earrings) and background warping.[9]

1.3 The Asymmetric Warfare of Detection

The relationship between deepfake generation and detection is often described as an "arms race" or a game of "cat and mouse".[10] Historically, detection methods have relied on identifying specific weaknesses in generation algorithms. For instance, early detectors focused on the lack of physiological signals, such as the pulse detectable in skin color variations (remote photoplethysmography) or irregular blinking patterns.[11]

However, these heuristic-based detectors suffer from a short shelf-life. Once a detector is published, its logic can be incorporated into the discriminator of a GAN, effectively training the generator to bypass that specific check. This adversarial dynamic necessitates a shift from "artifact-specific" detection (looking for a specific glitch) to "data-driven" detection (learning robust features from massive datasets) and "process-driven" detection (identifying the fundamental structural changes inherent to any manipulation).[12]

The Sentinel system represents this modern paradigm. It does not look for blinking. It looks for the statistical fingerprints left by the upsampling layers of a neural network and the blending artifacts introduced when a foreign object is inserted into a video frame. This approach aims for **generalization**—the ability to detect a deepfake created by a tool that did not exist when the detector was trained.[12]

2. Theoretical Framework: The Physics of Digital Forgery

2.1 The Spatial Domain vs. The Frequency Domain

To the human eye, a high-quality deepfake appears seamless. The skin texture looks realistic, the lighting matches the scene, and the expressions are convincing. However, the Sentinel system operates on the premise that digital images contain information invisible to biological vision systems.

In the **spatial domain** (the pixels we see), advanced GANs like StyleGAN2 are incredibly effective at fooling observers. They optimize the image to minimize the "perceptual loss," essentially mathematically guaranteeing that the image looks good. However, this optimization often neglects the **frequency domain**.

Research integrated into the Sentinel logic suggests that Convolutional Neural Networks (CNNs)—the backbone of most generators—have a systematic flaw: they struggle to replicate the high-frequency spectral distribution of natural images.[15] Natural images follow a specific statistical law regarding how energy is distributed across frequencies (typically a $1/f$ power law). When a GAN generates an image, it typically upsamples a low-resolution latent code. This upsampling process (often using transposed convolutions) leaves behind periodic "checkerboard" artifacts in the Fourier spectrum.

Sentinel's architecture, particularly the **EfficientNet-B7** backbone, is capable of learning these spectral anomalies. While the network operates on RGB pixels, the convolutional filters in its early layers can effectively act as band-pass filters, detecting the high-frequency noise residues that characterize synthetic content.[17]

2.2 The Blending Boundary Hypothesis (Face X-ray)

A critical theoretical pillar of the Sentinel system is the **Face X-ray** hypothesis. This theory posits that the vast majority of malicious deepfakes (specifically identity swaps) involve a two-step process:

1. **Generation:** Creating the fake face.
2. **Blending:** Merging this fake face into the original video frame.

Even if the generation step becomes perfect—indistinguishable from reality—the blending step creates a fundamental discontinuity. A real image is captured by a single sensor, meaning the noise profile (Photo Response Non-Uniformity or PRNU) and the resolution characteristics are consistent across the entire frame. A deepfake is a composite of two different sources: the background video and the generated face.[13]

The blending process, typically achieved using algorithms like Poisson Image Editing or Gaussian blurring, attempts to smooth the transition between these two sources. However, this smoothing itself alters the local image statistics, creating a detectable boundary. Sentinel's Face X-ray module is designed to detect this specific boundary artifact. By training the model to predict the *mask* of the manipulated region rather than just a binary "real/fake" label, the system learns to focus on the *edges* of the face, where the evidence of tampering is most likely to reside.[19]

2.3 Neuron Coverage and Activation Patterns

Beyond the visual and statistical artifacts, Sentinel's design is informed by the concept of **Neuron Coverage** or **FakeSpotter** methodology. The premise here is that real and fake faces trigger different activation patterns within a deep neural network.[21]

When a deep recognition network (like VGG-Face or ResNet) processes a real face, specific clusters of neurons fire to recognize familiar features (e.g., "nose," "eyebrow"). When it processes a fake face, even a very good one, the activation pathway is subtly different. The generated face might lack the "soul" or the intricate coherence of a real face, causing the network to rely on different, perhaps less robust, features for classification.

FakeSpotter monitors the layer-by-layer neuron behavior. It has been observed that fake faces tend to activate a different subset of neurons or result in lower overall activation consistency compared to real faces.[22] While Sentinel primarily uses EfficientNet for end-to-end classification, the high-dimensional feature vectors extracted by the model implicitly capture these activation discrepancies. The system's robustness is further enhanced by monitoring these internal states, providing a signal that is orthogonal to pixel-level artifacts.[21]

3. The Sentinel Architecture: Core Components

The Sentinel system is not a single algorithm but a composite architecture. It orchestrates data preprocessing, feature extraction, and classification into a unified pipeline. The core of this pipeline is the **EfficientNet-B7** convolutional neural network, trained using the **Noisy**

Student protocol.

3.1 EfficientNet-B7: The Backbone of Detection

The choice of EfficientNet-B7 is a calculated decision prioritizing accuracy and feature richness over raw speed. In the domain of deepfake detection, subtle artifacts—often spanning only a few pixels in a 1080p frame—are the only evidence of forgery. A lightweight model (like MobileNet) might compress these artifacts out of existence during the pooling layers.

Compound Scaling:
EfficientNet differs from traditional CNNs (like ResNet or VGG) in its scaling methodology. Traditional scaling involves arbitrarily increasing the depth (number of layers), width (number of channels), or resolution of the network. EfficientNet uses a compound coefficient (ϕ) to scale all three dimensions uniformly.23

- **Depth (α^ϕ):** A deeper network can capture more complex, hierarchical features.
- **Width (β^ϕ):** A wider network can capture more fine-grained features and is easier to train.
- **Resolution (γ^ϕ):** Higher resolution input images allow the network to see the high-frequency artifacts (checkerboard patterns) that are critical for forensics.

For Sentinel, EfficientNet-B7 (the largest variant in the original family) offers a massive receptive field and high input resolution (typically 600×600 pixels). This ensures that the minute discrepancies at the blending boundary or in the texture of the iris are preserved and analyzed.[24]

MBConv and Squeeze-and-Excitation:
The fundamental building block of EfficientNet is the Mobile Inverted Bottleneck Convolution (MBConv), enhanced with Squeeze-and-Excitation (SE) optimization.

- **MBConv:** Uses depthwise separable convolutions to reduce computational cost without sacrificing accuracy. This is crucial for B7, which is a heavy model; without MBConv, it would be too slow for any practical "real-time" application.[26]
- **Squeeze-and-Excitation:** This mechanism allows the network to perform dynamic channel-wise feature recalibration. In the context of Sentinel, the network learns to explicitly weight the channels that contain "forensic" information (e.g., noise residuals) higher than channels that contain semantic information (e.g., the identity of the person). This attention-like mechanism allows the model to focus on *validity* rather than *identity*.[23]

3.2 The "Noisy Student" Semi-Supervised Training Paradigm

Perhaps the most sophisticated aspect of Sentinel's algorithmic core is its training regime. Deep learning models are notoriously data-hungry, and while there are large deepfake datasets, they pale in comparison to the infinite supply of unlabeled video on the internet. Furthermore, models trained solely on clean, academic datasets often fail when deployed in the wild due to domain shift.

Sentinel addresses this via the **Noisy Student** algorithm, a semi-supervised learning method that iteratively improves the model's robustness.[27]

The Workflow:

1. **Teacher Training:** An EfficientNet-B7 is trained on the labeled **DFDC** dataset. This model learns the baseline artifacts of known deepfakes.
2. **Pseudo-Labeling:** This Teacher model is then used to infer labels on a massive unlabelled dataset (e.g., random YouTube videos or a large subset of the DFDC test set). It generates "soft" labels (e.g., "85% Fake").
3. **Student Training:** A new, equal or larger EfficientNet model (the Student) is trained on this combined dataset (labeled + pseudo-labeled). Crucially, the Student is trained with **noise**.
   - **Data Noise:** Heavy augmentation via Albumentations (compression, blur, cutout).
   - **Model Noise:** Dropout and Stochastic Depth (randomly skipping layers).
4. **Iterative Evolution:** The Student, forced to learn invariant features to match the Teacher's predictions despite the noise, becomes the new Teacher. The process repeats.

Forensic Implications:
The injection of noise is the key. If the Teacher says an image is a deepfake because of a specific pixel artifact, and the Student is shown a version of that image where that artifact is blurred out (via augmentation), the Student is forced to find other, more robust evidence (e.g., geometric inconsistency) to reach the same conclusion. This effectively simulates the challenging conditions of real-world detection, where videos are compressed and degraded.29

3.3 Face X-ray Integration

While EfficientNet provides the holistic classification, Sentinel integrates the logic of **Face X-ray** to handle generalizability. This component does not look for "artifacts" in the traditional sense but rather for the **blending boundary**.[13]

Self-Supervised Learning Mechanism:
The training for the Face X-ray module is unique because it does not require any deepfake images.

1. The system takes two real images, A and B.
2. It extracts the face from A and generates a random mask M.
3. It blends the face of A into image B using the mask M and standard blending techniques.
4. The network is trained to predict the mask M given the blended image.

Since the network learns to find the boundary M in these synthetic blends, and since real deepfakes *must* use similar blending to merge the generated face, the model learns a universal feature of face swapping. This makes Sentinel robust to "zero-day" deepfake generators; as long as they swap faces, they leave a blending trace that Face X-ray can detect.[13]

4. Data Engineering: The Foundation of Robustness

A machine learning model is only as good as the data it observes. Sentinel's engineering team has placed significant emphasis on constructing a training pipeline that mimics the variability of the real world.

4.1 The Deepfake Detection Challenge (DFDC) Dataset

The primary training corpus for Sentinel is the **Deepfake Detection Challenge (DFDC)** dataset. Released by Meta (Facebook) in collaboration with Microsoft and academic partners, this dataset represented a paradigm shift in forensic data.[32]

Dataset Characteristics and Strategic Value:

Feature	DFDC	Predecessors (e.g., FF++)	Sentinel's Benefit

Scale	~100,000+ clips	~5,000 clips	Prevents overfitting; exposes model to massive variance.
Subject Diversity	3,500+ paid actors (diverse age, race, gender)	Mostly white/young (YouTube scrapes)	mitigates **demographic bias**; ensures detection works across all skin tones.[33]
Generation Methods	8+ (DFAE, MM, NTH, FSGAN, StyleGAN)	4	Prevents the model from learning the quirks of a single generator.
Distractors	Audio swaps, Text-to-Speech (TTS)	None	Teaches the model to distinguish between visual fakes and legitimate videos with dubbed audio.[2]
Recording Environment	Unconstrained (dark rooms, outdoors, moving cameras)	Studio / News anchors	Prepares the model for "in-the-wild" footage found on social media.

Evaluation Metric: Log Loss:
The DFDC competition, and by extension Sentinel, optimizes for Log Loss (Logarithmic Loss) rather than simple accuracy. Log Loss heavily penalizes confident but wrong predictions. In a security context, this is vital. A detection system that flags a real video of a politician as fake with 99% confidence causes immense damage. Log Loss forces the model to be "calibrated"—to only output high confidence when the evidence is overwhelming.32

4.2 Albumentations: The Legacy of Data Science Bowl 2018

To implement the "Noisy Student" training, Sentinel relies on the **Albumentations** library. The selection of this specific library is not incidental; it traces its lineage to the winning solutions of the **2018 Data Science Bowl** (a nuclei segmentation competition).[34]

In the 2018 Data Science Bowl, the top teams realized that neural networks (like U-Net) performed poorly on unseen biological images because of variations in staining, lighting, and microscope quality. They solved this by applying aggressive, diverse data augmentations during training, forcing the model to learn the invariant shape of the nucleus rather than the color of the stain.[36]

Sentinel applies this exact philosophy to deepfakes. A deepfake should be detected regardless of the video compression (stain), the lighting (microscope), or the camera noise.

Key Albumentations Used in Sentinel:

1. **ImageCompression (JPEG/WebP):** Simulates the aggressive compression algorithms used by platforms like WhatsApp, Facebook, and TikTok. Deepfake artifacts are often high-frequency; compression removes high frequencies. By training on compressed images, Sentinel learns to find low-frequency structural artifacts that survive compression.[38]
2. **GaussNoise / ISONoise:** Simulates low-quality camera sensors (e.g., grainy webcam footage). This prevents the model from simply detecting the "clean" noise profile of a generator against a grainy background.[39]
3. **CoarseDropout / Cutout:** Randomly removes squares of the image. This forces the network to look at the *whole* face. If the network relies solely on "mismatched eyes," and the eyes are dropped out, it fails. CoarseDropout forces it to learn secondary features like ear rendering or mouth movements.[40]
4. **ShiftScaleRotate:** Ensures the model is invariant to the position of the face. Face detectors in the wild are imperfect; the face might be slightly off-center or tilted. The classifier must handle this.[35]

4.3 Preprocessing: Face Extraction and Alignment

Before the EfficientNet can see the data, it must be prepared.

- **Detector:** Sentinel likely uses a high-speed detector like **MTCNN** or **RetinaFace**.
- **Margin:** Crucially, the system extracts the face with a wide margin (e.g., 1.3x the bounding box). Why? Because the **Face X-ray** logic tells us the artifacts are at the *boundary*. A tight crop that cuts off the forehead or chin removes the blending zone, effectively destroying the evidence.
- **Alignment:** The faces are aligned (eyes made horizontal) to reduce geometric variance, making the job of the CNN easier.

5. Frequency Domain Analysis: The Invisible Evidence

While Sentinel primarily processes RGB pixels, its efficacy is deeply rooted in the physics of signal processing.

5.1 The Fourier Discrepancy

The **Fourier Transform** decomposes an image into its constituent sine and cosine waves. In a natural image, the amplitude of these waves decreases as the frequency increases (the image is mostly smooth with some edges).

Generative models, particularly those using **transposed convolutions** (the standard way to upsample a low-res latent vector to a high-res image), violate this natural law. They tend to insert excessive energy in the high-frequency bands. This manifests as a "grid" or "checkerboard" pattern in the frequency spectrum.[15]

- **Durall et al. (2019)** demonstrated that even when a GAN image looks perfect spatially, its spectral distribution is distorted. The high frequencies do not decay correctly.[11]
- **Sentinel's Leveraging:** By using high-resolution inputs (600×600), Sentinel ensures these high-frequency bands are present in the input signal. The initial convolutional layers of EfficientNet-B7, which act as edge detectors, pick up on these repetitive grid patterns. The **Squeeze-and-Excitation** blocks then amplify the feature maps corresponding to these unnatural frequencies, allowing the model to flag the image as fake even if the facial geometry is perfect.[16]

5.2 GAN Fingerprints and GANprintR

Every GAN architecture leaves a unique "fingerprint" based on its kernel size, number of layers, and upsampling method. This is similar to PRNU in camera forensics. However, relying on specific fingerprints is dangerous because they can be removed.

Research into **GANprintR** (GAN fingerprint Removal) shows that an autoencoder can be trained to suppress these specific high-frequency artifacts, effectively "washing" the deepfake to fool detectors that rely solely on fingerprints.[9]

Sentinel defends against this by **not** relying solely on high-frequency fingerprints. The multi-modal approach—combining the spatial structural checks of **Face X-ray** (which detects the blending, not the generation) with the semantic consistency checks of the **Noisy Student** training—ensures that even if the fingerprint is washed, the blending boundary or the lack of semantic coherence remains detectable.[44]

6. Performance Analysis and Benchmarking

The efficacy of the Sentinel system is quantified through rigorous benchmarking against standard datasets.

6.1 Quantitative Metrics

- **AUC (Area Under the Curve):** Sentinel-class systems (EfficientNet-B7 + Augmentation) typically achieve AUC scores of **0.95 - 0.98** on the DFDC test set.[25] This indicates a very high probability of ranking a randomly chosen fake video higher than a randomly chosen real one.
- **F1 Score:** On imbalanced datasets (where real videos far outnumber fakes), the F1 score is critical. Sentinel maintains a high F1 score (approx. **88.0%** [25]), indicating a good balance between precision (not flagging real videos) and recall (catching fake ones).
- **Log Loss:** As the primary optimization target, Sentinel aims for a Log Loss < 0.2. This ensures that the system is not just "guessing" right, but is statistically confident in its assertions.

6.2 Generalization and Robustness

The true test is **Cross-Dataset Generalization**.

- **Scenario:** Train on DFDC, Test on Celeb-DF.
- **Result:** Most classifiers see a massive performance drop (e.g., AUC drops from 0.99 to 0.60).
- **Sentinel Performance:** Due to Face X-ray (which detects the blending process common

to both datasets) and Albumentations (which simulates the quality differences), Sentinel maintains much higher generalization performance. Studies show Face X-ray alone can achieve **>95% AUC** on unseen datasets.[18]

Robustness to Perturbations:
When tested against videos with Gaussian blur or JPEG compression (Quality=40), standard Xception models often fail. The Noisy Student training ensures Sentinel retains accuracy, as it was explicitly trained to handle these degradations. EfficientNet-B7 has been shown to be significantly more robust to these common corruptions than ResNet architectures.24

6.3 Latency and Real-Time Viability

A significant tradeoff for using EfficientNet-B7 is computational cost. B7 has ~66 million parameters and a high FLOP count.

- **Inference Speed:** On a standard GPU (e.g., NVIDIA V100), B7 might run at ~20-30 FPS.
- **Optimization:** Sentinel likely uses **mixed-precision inference (FP16)** to double the throughput and **keyframe selection** (processing only 1 in every 5 frames) to achieve real-time performance on live streams. The system relies on the fact that deepfake artifacts are usually temporal; if they appear in frame t, they likely appear in $t+1$.[23]

7. Future Directions: The Evolving Battlefield

7.1 The Diffusion Threat

The emergence of **Diffusion Probabilistic Models** (e.g., Stable Diffusion, DALL-E 2 adapted for faces) represents a new challenge. These models generate images by iteratively denoising a random distribution.

- **Challenge:** They do not necessarily use the same upsampling layers as GANs, meaning the "checkerboard" frequency artifacts might be absent or different.
- **Sentinel Adaptation:** Future iterations of Sentinel will need to incorporate **Diffusion Fingerprint detection**, identifying the specific noise residuals characteristic of the diffusion process.[4]

7.2 Multi-Modal Analysis (Audio-Visual)

Deepfakes are increasingly multi-modal. A video might have a real face but a cloned voice (Audio Deepfake), or a swapped face with the original voice.

- **Lip-Sync Detection:** Sentinel is moving towards integrating **visual-audio dissonance detection**. By analyzing the phonemes (sound units) and visemes (mouth shapes), the system can detect if the mouth movement perfectly matches the speech. Models like **Wav2Lip** can generate lip-syncs, but they often leave temporal jitter that can be detected by recurrent networks (LSTMs) monitoring the mouth region.[46]

7.3 Edge Deployment and Privacy

There is a growing demand to run detection on-device (smartphones) to protect user privacy (no video uploaded to the cloud).

- **Distillation:** Techniques like **Knowledge Distillation** can be used to compress the massive EfficientNet-B7 into a smaller **EfficientNet-B0** or **MobileNet**. The B7 acts as the "Teacher" and the mobile model as the "Student," learning to mimic the B7's output with a fraction of the parameters.[25]

8. Conclusion

The Sentinel Real-Time Deepfake Detection System stands as a sophisticated technological bulwark against the erosion of truth in digital media. Its architecture is a testament to the complexity of modern forensics: it is not enough to simply "look" at an image. One must understand the generative process, the statistical properties of digital signals, and the adversarial nature of the domain.

By fusing the raw feature-extraction power of **EfficientNet-B7** with the structural insight of **Face X-ray** and the robust training methodology of **Noisy Student**, Sentinel achieves a level of generalization that represents the state-of-the-art. It acknowledges that the "perfect" deepfake is a moving target, and thus focuses on the fundamental invariants of forgery—blending boundaries and spectral anomalies—rather than transient visual flaws.

However, the system is not a silver bullet. The rapid advance of diffusion models and holistic neural rendering requires constant evolution. As Sentinel evolves, it will likely transition from a purely visual inspector to a holistic multi-modal analyst, ensuring that as the line between real and synthetic blurs, we retain the tools to illuminate the difference.

Works cited

1. (PDF) DeepFakes and Beyond: A Survey of Face Manipulation and Fake Detection, accessed November 19, 2025, https://www.researchgate.net/publication/338355353_DeepFakes_and_Beyond_A_Survey_of_Face_Manipulation_and_Fake_Detection
2. DeepFake Detection Challenge - IQT, accessed November 19, 2025, https://www.iqt.org/library/deepfake-detection-challenge
3. Deep Fake Image Detection Based on Pairwise Learning - MDPI, accessed November 19, 2025, https://www.mdpi.com/2076-3417/10/1/370
4. DiffusionFake: Enhancing Generalization in Deepfake Detection via Guided Stable Diffusion - arXiv, accessed November 19, 2025, https://arxiv.org/html/2410.04372v1
5. UGAD: Universal Generative AI Detector utilizing Frequency Fingerprints - arXiv, accessed November 19, 2025, https://arxiv.org/html/2409.07913v1
6. "Deepfakes and beyond: A Survey of face manipulation and fake detection - UAM, accessed November 19, 2025, https://repositorio.uam.es/bitstream/10486/702969/1/tolosana_deepfakes_InfFus_2020_pre.pdf
7. Deepfakes Generation and Detection: A Short Survey - PMC - PubMed Central, accessed November 19, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC9863015/
8. Deepfakes Generation and Detection: State-of-the-art, open challenges, countermeasures, and way forward - SciSpace, accessed November 19, 2025, https://scispace.com/pdf/deepfakes-generation-and-detection-state-of-the-art-open-fgrtufi4.pdf
9. GANprintR: Improved Fakes and Evaluation of the State of the Art in Face Manipulation Detection - Julian FIERREZ, accessed November 19, 2025, https://biometrics.eps.uam.es/fierrez/files/2020_JSTSP_GANprintR_Neves.pdf
10. The Tug-of-War Between Deepfake Generation and Detection - arXiv, accessed November 19, 2025, https://arxiv.org/html/2407.06174v4
11. Amplifying attention to artifacts increases deepfake detection by humans and machines, accessed November 19, 2025, https://arxiv.org/html/2206.00535v4
12. DEEPFAKE DETECTION WITH CONTRASTIVE LEARN- ING IN CURVED SPACES - OpenReview, accessed November 19, 2025, https://openreview.net/pdf?id=G4D6jClNFI
13. [1912.13458] Face X-ray for More General Face Forgery Detection - arXiv, accessed November 19, 2025, https://arxiv.org/abs/1912.13458
14. Veritas: Generalizable Deepfake Detection via Pattern-Aware Reasoning - arXiv,

accessed November 19, 2025, https://arxiv.org/html/2508.21048v1
15. Leveraging Frequency Analysis for Deep Fake Image Recognition - Proceedings of Machine Learning Research, accessed November 19, 2025, https://proceedings.mlr.press/v119/frank20a/frank20a.pdf
16. A Closer Look at Fourier Spectrum Discrepancies for CNN-generated Images Detection, accessed November 19, 2025, https://www.researchgate.net/publication/355861820_A_Closer_Look_at_Fourier_Spectrum_Discrepancies_for_CNN-generated_Images_Detection
17. A Closer Look at Fourier Spectrum Discrepancies for CNN-Generated Images Detection - CVF Open Access, accessed November 19, 2025, https://openaccess.thecvf.com/content/CVPR2021/papers/Chandrasegaran_A_Closer_Look_at_Fourier_Spectrum_Discrepancies_for_CNN-Generated_Images_CVPR_2021_paper.pdf
18. Face X-Ray for More General Face Forgery Detection - CVF Open Access, accessed November 19, 2025, https://openaccess.thecvf.com/content_CVPR_2020/papers/Li_Face_X-Ray_for_More_General_Face_Forgery_Detection_CVPR_2020_paper.pdf
19. Deepfake Detection via 3D Face Reconstruction-Based Image, accessed November 19, 2025, https://www.computer.org/csdl/proceedings-article/hpcc/2025/687400a230/2bggmEnqTkI
20. Vulnerability-Aware Spatio-Temporal Learning for ... - OpenReview, accessed November 19, 2025, https://openreview.net/pdf/d218f62b6d19ef551d0ee1e1f0295594ad561298.pdf
21. FakeSpotter: A Simple yet Robust Baseline for Spotting AI-Synthesized Fake Faces - arXiv, accessed November 19, 2025, https://arxiv.org/abs/1909.06122
22. FakeSpotter: A simple yet robust baseline for spotting AI-synthesized fake faces, accessed November 19, 2025, https://kyushu-u.elsevierpure.com/en/publications/fakespotter-a-simple-yet-robust-baseline-for-spotting-ai-synthesi
23. EfficientNet in deepfake detection - Antispoofing Wiki, accessed November 19, 2025, https://antispoofing.org/efficientnet-for-deepfake-detection-architecture-purpose-and-practical-application/
24. Deepfake Detection Using EfficientNetB7: Efficacy, Efficiency, and Adaptability, accessed November 19, 2025, https://ijisae.org/index.php/IJISAE/article/view/5556
25. (PDF) Combining EfficientNet and Vision Transformers for Video Deepfake Detection, accessed November 19, 2025, https://www.researchgate.net/publication/353042472_Combining_EfficientNet_and_Vision_Transformers_for_Video_Deepfake_Detection
26. Enhancement of Deepfake Detection Framework Integrating EfficientNetB0, Graph Attention Networks (GAT) and Gated Recurrent Units (GRU) - ResearchGate, accessed November 19, 2025, https://www.researchgate.net/publication/392171020_Enhancement_of_Deepfake_Detection_Framework_Integrating_EfficientNetB0_Graph_Attention_Networks_

GAT_and_Gated_Recurrent_Units_GRU
27. Self-training with Noisy Student improves ImageNet classification | Request PDF, accessed November 19, 2025, https://www.researchgate.net/publication/337183832_Self-training_with_Noisy_Student_improves_ImageNet_classification
28. Self-Training With Noisy Student Improves ImageNet Classification - CVF Open Access, accessed November 19, 2025, https://openaccess.thecvf.com/content_CVPR_2020/papers/Xie_Self-Training_With_Noisy_Student_Improves_ImageNet_Classification_CVPR_2020_paper.pdf
29. Self-Training With Noisy Student Improves ImageNet Classification - IEEE Xplore, accessed November 19, 2025, https://ieeexplore.ieee.org/document/9156610/
30. [1911.04252] Self-training with Noisy Student improves ImageNet classification - arXiv, accessed November 19, 2025, https://arxiv.org/abs/1911.04252
31. arXiv:1912.13458v2 [cs.CV] 19 Apr 2020, accessed November 19, 2025, https://arxiv.org/pdf/1912.13458
32. arXiv:1910.08854v2 [cs.CV] 23 Oct 2019, accessed November 19, 2025, https://arxiv.org/pdf/1910.08854
33. Deepfake Detection Challenge Results: An open initiative to advance AI - AI at Meta, accessed November 19, 2025, https://ai.meta.com/blog/deepfake-detection-challenge-results-an-open-initiative-to-advance-ai/
34. [ods.ai] topcoders, 1st place solution - Kaggle, accessed November 19, 2025, https://www.kaggle.com/competitions/data-science-bowl-2018/writeups/ods-ai-topcoders-ods-ai-topcoders-1st-place-soluti
35. Albumentations: Fast and Flexible Image Augmentations - MDPI, accessed November 19, 2025, https://www.mdpi.com/2078-2489/11/2/125
36. mistrzuniu1/Data-Science-Bowl-2018---kaggle-competition - GitHub, accessed November 19, 2025, https://github.com/mistrzuniu1/Data-Science-Bowl-2018---kaggle-competition
37. 2018 Data Science Bowl | Kaggle, accessed November 19, 2025, https://www.kaggle.com/competitions/data-science-bowl-2018
38. Documentation - Albumentations, accessed November 19, 2025, https://albumentations.ai/docs/
39. Albumentations: fast and flexible image augmentations, accessed November 19, 2025, https://albumentations.ai/
40. albumentations-team/albumentations: Fast and flexible image augmentation library. Paper about the library: https://www.mdpi.com/2078-2489/11/2/125 - GitHub, accessed November 19, 2025, https://github.com/albumentations-team/albumentations
41. A Closer Look at Fourier Spectrum Discrepancies for CNN-generated Images Detection - Keshigeyan Chandrasegaran, accessed November 19, 2025, https://keshik6.github.io/pdfs/nvidia_talk_cvpr21_oral_slides.pdf
42. (PDF) GANprintR: Improved Fakes and Evaluation of the State of the Art in Face Manipulation Detection - ResearchGate, accessed November 19, 2025, https://www.researchgate.net/publication/337241864_GANprintR_Improved_Fake

s_and_Evaluation_of_the_State_of_the_Art_in_Face_Manipulation_Detection
43. [1911.05351] GANprintR: Improved Fakes and Evaluation of the State of the Art in Face Manipulation Detection - arXiv, accessed November 19, 2025, https://arxiv.org/abs/1911.05351
44. Evaluating the Effectiveness of a Gan Fingerprint Removal Approach in Fooling Deepfake Face Detection - Journal of Internet Services and Information Security, accessed November 19, 2025, https://jisis.org/wp-content/uploads/2024/03/2024.I1.006.pdf
45. Comparison of ResNet, EfficientNet, and Xception architectures for deepfake detection - CEUR-WS.org, accessed November 19, 2025, https://ceur-ws.org/Vol-3899/paper3.pdf
46. FakeHunter: Multimodal Step-by-Step Reasoning for Explainable Video Forensics - arXiv, accessed November 19, 2025, https://arxiv.org/html/2508.14581v1
47. Deepfake Detection using TCN and EfficientNet- B3 - IJSAT, accessed November 19, 2025, https://www.ijsat.org/papers/2025/2/3646.pdf